

Vyra Landscape

The front door to Vyra's canonical docs, and a stakeholder-facing read on its own: the vision, the operating model, and how a regulation becomes provable assurance. It carries **no build-status** — for what's live vs. planned, see `vyra-tracker.md`. Also: `vyra-graph-spine.md` for the ground-truth schema, `vyra-architecture.md` for the software layers, `vyra-implementation-plan.md` for phase sequencing.

What Vyra Is

Vyra is an Agentic Risk & Compliance Infrastructure platform. AI agents continuously transform regulations into operational assurance by reasoning over a shared enterprise graph — the **Compliance Digital Twin**.

The operating loop, end to end:

Regulations → Requirements → Execution → Operations → Signals

That loop runs across five graph domains (full detail in `vyra-graph-spine.md`):

Graph	Question
Knowledge	What must be done?
Execution	What are we doing?
Operational	What is happening?
Intelligence	What do we understand?
Assurance	What can we prove?

Operating principles (the subset that governs day-to-day design decisions, not the full pitch):

- Agents collaborate through the graph, not through messaging — the graph is the operating system, not the product.
- Every compliance outcome must be traceable, forward and reverse (mechanics: `vyra-graph-spine.md`'s Graph Traversal Patterns).
- Every agent follows the same lifecycle: **Observe → Interpret → Reason → Act → Verify → Learn**.
- Not every action gets full autonomy. **Autonomy Levels 0–4** — Human Driven → Agent Recommends → Agent Assisted → Agent Executed → Fully

Autonomous. Default is **Level 1 (Agent Recommends, Human Approves)** unless a specific workflow is explicitly elevated.

Vyra — Compliance. Handled.

The Operating Model — 7-Layer JTBD

Sourced from the GRC Operating Model bow-tie diagram. Each layer names a **persona** (who does the job today, and which agent family eventually automates it) and its **capabilities** — the jobs-to-be-done Vyra exists to serve.

This table is the operating model, not a status report. For where each capability stands today (live / partial / gap) and what backs or blocks it, see `vyra-tracker.md`.

L	Layer	Persona → Agent Family	Capabilities
L1	Knowledge	Catalog Admin → Regulatory Intelligence Agents	Regulations · Standards · Contracts · SOPs
L2	Interpret	Ops Admin → Applicability Intelligence Agents	Applicability Scoping · Obligation Linkage
L3	Planning	Planner → Control Intelligence Agents	52-Week Calendar · Location + Role Assignment
L4	CTN Knowledge Graph Spine — Capture · Review	<i>(the graph itself — every agent's shared memory)</i>	The shared substrate all other layers read and write; "Review" = the Autonomy Level 1 human-approval gate
L5	Oversight	Ops Supervisor → Signal Intelligence Agents	Deviation Alerts · Escalation Paths
L6	Assurance	Compliance Mgmt → Assurance Agents	Coverage Scoring · Audit-Ready Export
L7	Risk	Risk Manager → Risk Intelligence	Residual Risk Score · Scenario Simulation

Where to go next

- **Current build status** (what's live vs. partial vs. gap, per JTBD layer) → [vyra-tracker.md](#)
- **Schema, relationships, live/dormant status, Cypher patterns** → [vyra-graph-spine.md](#)
- **Software layers & components** (how the API, agents, ingestion, and UI wrap the graph) → [vyra-architecture.md](#)
- **Sequencing, what's done, what's next, verification approach** → [vyra-implementation-plan.md](#)
- **Historical material** (original blueprint narrative, superseded graph/roadmap docs, drift-analysis working notes, synthetic seed data) → [__ref/](#)